

A complaint should not disappear between desks.

Axiom turns one customer email and PDF into a visible, governed outcome — in minutes, not hours.

01

**CUSTOMER
WRITES**

02

**BANK
PROVES**

The customer waits. The bank loses the thread.

90 min

baseline handling time
per complaint

11%

regulatory deadlines
missed

1.8M

retail and SME
customers affected

The real failure

**Evidence,
ownership
and
deadlines
drift
between
teams.**

So speed alone is not enough. Every handoff must stay understandable, safe and provable.

One complaint. One visible team. One signed outcome.

01

Customer writes

Their own words and PDF become the case input.



02

A visible team checks

Every handoff shows who acted and what they found.



03

Rules protect the outcome

AI may suggest; bank rules decide money, deadlines and disclosures.

Nothing is pre-written except the fictional identity.

VISITOR INPUT

Type any complaint

Use your own subject and words.

Upload your PDF

The amount, merchant and evidence filename follow what you submit.

Demo safety

Only the customer name, email and account are fictional.

Write your own complaint
Your wording, your evidence

Use the sample
One fixed sanitized email

RFC822 / YOU WRITE IT

SYNTHETIC

File it as

Ahmad bin Ismail · *****6890 · savings

This choice supplies only a safe fictional identity and matching ledger account. Your complaint, amount, merchant and PDF are not prefilled; they become the fresh case.

Sender name

Ahmad bin Ismail

Disputed amount (RM)

2450

Subject

Card charge I never made

24 / 168 · MINIMUM 4

What happened?

A card payment left my account last night and I never approved it. Please reverse it.

85 / 4000 · MINIMUM 20

Do not paste a real account number or NRIC. The runner refuses them outright rather than quietly deleting them, so nothing teaches you that real data is accepted here.

Merchant

TECHWORLD KL

Evidence PDF (optional)

Choose File my-evidence.pdf

A digital PDF is read as text; a scan is transcribed by the vision model. Either way the transcription is screened again before any reasoning step, because a document can carry an instruction as easily as an email can.

FROM

Ahmad bin Ismail <ahmad.ismail@example.my>

ACCOUNT

*****6890

Run My Complaint Live

INPUT / YOUR WORDS

The screen gives the visitor their input receipt.

INPUT RECEIPT / WRITTEN BY THIS VISITOR

This result starts with what you submitted.

The fictional identity only keeps real customer data out of the public demo. The complaint wording, amount, merchant and PDF below are the inputs the pipeline used.

SUBJECT	CLAIM
Card charge I never made	RM 2,450.00 · TECHWORLD KL
EVIDENCE	WORDS RECEIVED
my-evidence.pdf · Digital text read directly	80 characters · SHA-256 aaaaaaaaaaaaaa...

INPUT RECEIPT

Subject
exact visitor text

Claim
amount + merchant

Evidence
chosen PDF + read method

Body
length + unique fingerprint

Then the result shows the final state, event chain and balanced journal.

The most important feature is a safe ‘no’.


AI interprets

the complaint, category and customer explanation

Rules decide

money movement, deadlines, authority and mandatory disclosures

BLOCKED BEFORE AI



Start where the value is high and the path is real.

90 min

case-study baseline



< 30 sec

latest verified public run

57%

of stated volume in the first rollout

NOW

Synthetic pilot

Deployed app, real models, bank-tool simulators and database proof.

NEXT

Controlled rollout

Two high-volume dispute types with staff review.

THEN

Bank integration

Bank credentials, compliance sign-off and security approval.

Give us one complaint.

We'll give every customer, operator and auditor the same provable story.

casezero-alpha.vercel.app/live

Visitor-written complaint · Visitor-uploaded PDF · Fictional identity only

[Q&A →](#)

Who benefits — and why should they care?

Who is the primary user?

A small complaints team that needs one clear view of the customer, evidence, deadline and next action.

What changes for the customer?

They get a faster, clearer answer and a traceable explanation instead of repeating their story across teams.

What changes for the bank?

Less manual chasing, fewer missed deadlines and one auditable record from intake to final response.

How do you measure value?

Handling time, deadline misses, first-time resolution, safe-refusal rate and staff time saved—reported separately as measured, baseline or assumption.

Where is AI meaningful — and what is new?

Where is AI used?

AI reads messy complaint language, classifies the issue, compares evidence and drafts a plain-language customer explanation.

What does AI never decide?

It cannot move money, change a deadline, skip a required disclosure or grant itself authority.

What makes CaseZero different?

Governance is visible product behavior: every handoff leaves a receipt, every write needs a scoped ticket, and safe refusal is treated as success.

Why six specialists?

Each specialist has one narrow job. That makes handoffs understandable and lets the bank improve or replace one part without hiding the whole decision.

How do you keep people and money safe?

What happens if the AI is uncertain?

The case stops for human review. It does not guess its way into a financial action.

How do you handle prompt injection?

Suspicious instructions inside email or PDF evidence are blocked before any model is called.

How is personal data protected?

The public demo uses fictional identities. In deployment, identifiers are encrypted, redacted in views and restricted by role.

Can staff override the system?

Authorized staff can review and resolve exceptions, but the action and reason remain visible in the audit trail.

Can a bank actually adopt this?

Is this only a prototype?

The dashboard and API are deployed today with real model calls, a real database and controlled bank-tool simulations.

What still needs the bank?

Core-system credentials, complaint-mailbox access, staff identities, compliance sign-off and production security approval.

How would you start?

Begin with synthetic cases, then a staff-reviewed pilot for two high-volume dispute types before allowing controlled financial writes.

What if a bank already has a case system?

CaseZero can sit above existing systems as the governed workflow and proof layer instead of replacing everything at once.

What exactly can a visitor test?

Is the complaint prefilled?

No. Only a fictional identity is provided. The subject, complaint words, amount, merchant and uploaded PDF come from the visitor.

Does the PDF really matter?

Yes. Its filename and read method appear in the input receipt, and its evidence is compared with the visitor's claim.

What if the PDF is scanned?

The system can read scanned pages visually; if the evidence is unclear or inconsistent, it stops for review.

Is it accessible on mobile?

The live workspace is keyboard-operable, has visible focus, reduced-motion support, readable states and no horizontal overflow at 390px.

How do you know the result is real?

What proof does one run leave?

A visitor-input receipt, specialist handoff receipts, a final verdict, a balanced journal and a linked event history.

Can a model fake a successful result?

No. The final state and money movement come from server-side rules and bank-tool responses, not from model wording.

How is it tested?

The release gate runs 500+ backend tests, 29 browser journeys, type checks, a production build and dependency security checks.

What should we verify live?

Use a unique subject and PDF, then confirm the input receipt, final state, valid event links, journal balance and safe-refusal path.