

Bank complaints resolved in minutes — governed end to end.

One customer complaint becomes one visible decision that the operator, customer and auditor can all understand.

ONE
COMPLAINT


SIX
SPECIALISTS

+ one human-visible gate

A complaint takes 90 minutes because the evidence keeps changing hands.

90
minutes

Case-study baseline for one complaint moving across multiple departments.




The failure is not one slow person. It is ownership drift between desks.

When ownership drifts, customer delay becomes regulatory risk.


11%
of regulatory deadlines
missed today

1.8M
retail and SME
customers in the case study



A faster answer only counts if the bank can explain who acted, what evidence they used, and why money moved.

CaseZero turns one complaint into one visible, governed outcome.



The product is the chain of responsibility—not a chatbot added to a dashboard.

AI can read the mess. It should not own the irreversible decision.

01

Narrow roles

One responsibility per agent
makes mistakes visible.

02

Deterministic authority

Code—not model wording—owns
deadlines and money.

03

Human exception lane

Uncertainty becomes a review
task, never a guess.

Why now: modern models can interpret messy language and documents; governed software can keep their authority bounded.

Six specialists divide the work so no single brain owns the case.

01

GEMINI

Intake

Reads email + PDF; blocks harmful instructions.

02

GEMINI

Classifier

Chooses 1 of 7 categories; reports confidence.

03

NO MODEL

Verifier

Compares the claim with CRM and core banking.

04

NO MODEL

Resolver

Posts only with a signed kernel ticket.

05

GROQ

Communicator

Drafts the customer answer; rules insert legal text.

06


NO MODEL

Supervisor

Watches every deadline and escalates to people.

Compliance Kernel sits between them all: no model, no persuasion, only policy.

Models propose. The kernel decides. A human handles uncertainty.



Human enters for low confidence, incomplete or contradictory bank evidence, approval limits, or dual control.

The complaint and PDF come from the visitor—not from seeded text.

Only the identity is fictional

The visitor controls the subject, complaint words, amount, merchant and evidence PDF.

LIVE DEMO CUE

Type → attach → run

The first result block returns the exact subject, claim, PDF filename, read method and body fingerprint.

AXIOM / LIVE CASE 01

Watch one complaint become proof.

Write the complaint yourself, attach your own PDF, and watch each stage land as its evidence is written. No dashboard tour, no prefilled result — the pipeline underneath is the deployed one.

Write your own complaint
Your wording, your evidence

Use the sample
One fixed sanitized email

RFC822 / YOU WRITE ITSYNTHETIC

File it as

Ahmad bin Ismail · *****6890 · savings

This choice supplies only a safe fictional identity and matching ledger account. Your complaint, amount, merchant and PDF are not prefilled; they become the fresh case.

Sender name

Ahmad bin Ismail

Disputed amount (RM)

180.50

Subject

Aurora Books charge I did not authorise

39 / 160 · MINIMUM 4

What happened?

A midnight charge from Aurora Books appeared on my card and I did not authorise it. Please inspect my uploaded statement and reverse this payment.


146 / 4000 · MINIMUM 20

Do not paste a real account number or NRIC. The runner refuses them outright rather than quietly deleting them, so nothing teaches you that real data is accepted here.

Merchant

Evidence PDF (optional)

Each agent speaks only from evidence already written to the chain.



The minutes

EVERY LINE DERIVES FROM A HASHED EVENT. CLICK ONE TO CHECK IT.

SEQ 2

INTAKE → CLASSIFIER · CHECK FOR HARMFUL CONTENT

Clean. English complaint, claiming RM2,450.00, against TECHWORLD KL. Account number encrypted before anything read it.

+1.0s

SEQ 4

CLASSIFIER → COMPLIANCE KERNEL · IDENTIFY THE COMPLAINT TYPE

This is an unauthorised transaction. I am 97% confident.

+2.0s

SEQ 5

COMPLIANCE KERNEL → VERIFIER · SET PRIORITY AND DEADLINE

High priority, 5 working days under the BNM window, due 2026-08-11. Unauthorised transactions are treated as high urgency.

+3.0s

SEQ 7

VERIFIER → COMPLIANCE KERNEL · CHECK AGAINST BANK RECORDS

Core banking confirms the claim. 3 checks against the ledger. The transaction exists, the amount matches, and the card was not present.

+4.0s

VERIFICATION	PASS
EVIDENCE CHECKS	3
BANK TOOL CALLS	2
MODEL CALLS	none — this seat compares records, it does not reason
VERIFICATION_COMPLETED · SEQ 7 ·	e414ce2b6786843f0b...

SEQ 9

COMPLIANCE KERNEL → RESOLVER · DECIDE: AUTO-RESOLVE OR HUMAN REVIEW

Approved for automatic resolution. PASS evidence with the amount inside the auto-approval limit.

+5.0s

SEQ 10

RESOLVER → COMMUNICATOR · MOVE THE MONEY

RM2,450.00 reversal posted. Debit SUSPENSE:FRAUD, credit *****6890. Books balance.

+6.0s

SEQ 13

COMMUNICATOR → CUSTOMER · SEND RESPONSE TO CUSTOMER

Letter written and sent. Compliance lint passed with no repairs. The regulatory wording is inserted by the kernel, not by me.

+7.0s

Click any spoken line → see event number, facts and hash.

The human is not a fallback screen. They are the decision owner for exceptions.

The system assembles

evidence · category ·
confidence · deadline · bank
result

The person decides

Approve
Reject
Request information

Every choice re-enters the same kernel
gates and remains on the audit trail.

The screenshot displays the Axiom decision interface. On the left is a sidebar with navigation links: OPERATIONS (01 Today, 02 Dashboard, 03 Review queue), COMPLIANCE (04 Audit trail, 05 Blocked intake, 06 Fraud patterns), SETUP (07 Rules, 08 Operators, 09 Settings), and Synthetic rehearsal register. The main content area is titled 'INVESTIGATOR CO-PILOT / TARGET UNDER 60 SECONDS' and features the heading 'Decide with the evidence assembled.' Below this is a table with columns: CASE, CLAIM, VERIFICATION, URGENCY, and SLA DUE. The table contains one row for case MYB-2026-000014 with a claim of RM 486.20 and a verification status of PASS. To the right of the table is a 'DECISION REGISTER' section for the same case, showing 'Claimed RM 486.20', 'Category billing error', and 'Confidence 82%'. Below the register is a 'Decision note' text area and two buttons: 'Approve resolution' and 'Reject claim'. At the bottom right, there is a small 'Axiom' logo.

CASE	CLAIM	VERIFICATION	URGENCY	SLA DUE
MYB-2026-000014	RM 486.20 Duplicate annual service fee needs a human approval.	PASS	Medium	27/08/2026

DECISION REGISTER
MYB-2026-000014

Claimed RM 486.20

Category billing error

Confidence 82%


Decision note

Approve resolution

Reject claim

Axiom

The business case is measurable staff time plus a controlled path to rollout.



57%
unauthorised transactions + billing errors
in the case study

- NOW
Synthetic pilot
- NEXT
Staff-reviewed rollout
- THEN
Bank integration

One fresh visitor-written case proves the full chain works.

20.1s

fresh live case

14

valid event links

3 + 3

model + bank calls

RM180.50


balanced posting

VISITOR INPUT RECEIPT

Their subject · their claim · their PDF · exact body fingerprint

VERIFIED LIVE

Approve a staff-reviewed pilot for unauthorised transactions and billing errors.




Models propose. Rules decide. Humans own exceptions.

casezero-alpha.vercel.app/live

READY FOR A CONTROLLED PILOT

Natural language helps the human operate—without inventing authority.



- Axiom may
 - summarise
 - navigate
 - prepare an action
- Axiom may not
 - grant itself a role
 - bypass confirmation
 - post without a ticket

Likely questions about the six agents and their brains

Why six agents instead of one?

Each agent owns one narrow responsibility. That makes failures inspectable and prevents one model from controlling the full case.

Which agents use AI?

Intake and Classifier use Gemini; Communicator uses Groq. Verifier, Resolver and Supervisor deliberately make no model call.

Can you change a model?

Yes. Brain routing is separate from policy and authority, so a model can be evaluated or replaced without changing the financial gate.

What does the Compliance Kernel do?

It calculates deadlines, checks confidence and evidence, enforces limits, mints signed tickets and inserts mandatory customer wording.

Likely questions about people, safety and ethics

Exactly when does a human enter?

Low confidence, incomplete or contradictory evidence, an amount above the automatic limit, dual control, or any explicit refusal.

Can a human override everything?

An authorized person can decide an exception, but the reason is recorded and the action must still pass the same kernel and role checks.

What if the PDF contains an attack?

Suspicious instructions are blocked before the first model call. The case is quarantined with an explanation instead of silently continuing.

What if the AI is wrong?

Its output is only a proposal. Bank evidence and deterministic rules can contradict it, and uncertainty routes to a person rather than money movement.

Likely questions about value, rollout and adoption

Who pays and why?

A bank complaints operation pays to reduce handling time and deadline risk while preserving review, audit and existing control ownership.

Is 20 seconds realistic?

It is the measured synthetic production run time, not a promise for every real case. Network, evidence and human review can make cases longer.

How would rollout start?

Synthetic cases first, then a staff-reviewed pilot for two high-volume categories, then controlled integration after compliance and security approval.

What still belongs to the bank?

Core and CRM credentials, mailbox transport, operator identities, approval limits, policy ownership, compliance sign-off and production security.

Likely questions about the demo and technical proof

Is the complaint prefilled?

No. Only the demo identity is fictional. Subject, words, amount, merchant and PDF come from the visitor and appear in the input receipt.

How do you prove one result?

The result includes the visitor-input receipt, agent handoffs, model and bank-tool receipts, balanced journal and a recomputed event chain.

How is it tested?

The release gate passes 504 backend tests, 29 browser journeys, type checking, production build, slide/PDF rendering and dependency audit.

Is the interface accessible?

The live path supports keyboard use, visible focus, reduced motion, readable refusal states and zero horizontal overflow at 390 pixels.